

AI/ML Onboarding & Agent OS Plan

Day-by-day playbook (product work + agent setup merged)

For joining a multi-company ML platform (lead / churn / NBA-like):
build eval & degradation foundations while configuring agents for leverage.

Week 1 constraint: meetings-heavy — about **2 hours/day** for deep work.

Updated 2026-07-30

Contents

1. How to use this document
2. Context, strategy & principles
3. Week overview at a glance
4. Days 0–5 (Week 1) — meetings + 2h/day map
5. Days 6–12 — eval harness is the product
6. Days 13–18 — health wedge + promotion checklist
7. Days 19–30 — degradation, ownership, freeze
8. Canonical workflows & Agent OS reference
9. Repo foundation plan (team epic paste)
10. Checklists, red flags, AGENTS.md skeleton

1. How to use this document

Sections 4–7 are the **calendar**: each day merges *what to learn/ship* and *what agent tooling to stand up*. Do not maintain two separate plans.

- **Week 1:** ~2h deep work/day. Prefer listening, mapping, and tiny vault notes — no harness coding yet.
- **Weeks 2–4:** normal focus blocks; building the proof layer is the job.
- **Section 9:** team-facing epic text you can paste into the repo later.
- Metaphor: **you = eng manager; agents = team**. You own metrics, golden cases, promotion judgment, merges.

2. Context, strategy & principles

Project reality

- Many models per company (lead, churn, NBA, ...); companies roll out as data gets fixed.
- No eval harness/dashboard yet; no easy API/business metrics view; no clear degradation signal.

Fix data to unlock companies, but unlock only as fast as you can prove models still work — eval + health + promotion gates before scaling to hundreds of models.

Deliberate-practice split (from week 2 onward)

Bucket	%	Here means
Ship proof loop	40%	Eval runner, baselines, checklist, health
Evals & cases	25%	Golden set from real failures; worst-slice
Domain / data	20%	Why companies fail promotion; offline/online gaps
Agentic craft	10%	Spec → agent implements → you verify
Reading	5%	Only if it unblocks today's artifact

Trust rules (every day)

- Fresh-context review — never trust the authoring session to certify itself.
- Escalate ambiguous / product / threshold / promotion decisions to you.
- End-to-end evidence (report.json, health output) — not only unit tests.
- One task → one branch → one PR; keep work auditable.

Week 1 operating rule

With ~2h/day, your only jobs are: **boot once, map the loop, capture folklore, pick a pilot hypothesis, seed a tiny knowledge vault**. Do **not** start the eval runner, multi-agent loops, or dashboard work in week 1. Meetings are the data source — turn them into notes agents will read later.

3. Week overview at a glance

Days	Deep-work budget	Merged theme	Must-have artifacts
0–5 (Wk1)	~2h/day	Meetings + map + vault seed	System card, inventory draft, pain list, AGENTS.md stub
6–12	full blocks	Eval harness v0	Metric spec, golden set, report.json, fault inject
8–14	parallel	Serving health wedge	health_check + 5-min runbook
13–18	full blocks	Promotion checklist	Data vs model readiness dry-run
19–24	full blocks	Degradation path	CI/schedule gate + "how we know healthy"
25–30	full blocks	Own + freeze Agent OS	Ownership note, 60d plan, agent-os.md

Build order for Agent OS: vault + AGENTS.md → harness stubs → commands → skills → loops → hooks. Never invent 10 skills before one harness command exists.

4. Days 0–5 (Week 1) — meetings + ~2h/day

Assumption: calendar is full. Protect ~2 focused hours. Prefer async notes over coding. Every meeting should feed the vault (who said what about failures, data blockers, model types).

Day 0 — Arrive / access

Time budget: ~2h | *Mode:* meetings-heavy — listen, map, don't build

Goal: Repo boots; you know where docs and CI live.

Do today (product + agent OS merged):

- Get accounts, clone, install, run existing test/CI once (even if partial).
- Skim README / architecture only — no deep-doc binge.
- In meetings: ask who owns training, serving, data promotion, on-call.
- Start a private note: “how people notice breakage today” (Slack? never?).

Agent OS touch (keep tiny):

- Create empty stubs only if cheap: AGENTS.md (5 lines), docs/onboarding/system-card.md template, knowledge/INDEX.md with 4 domain names.
- Do **not** add skills, hooks, overnight loops, or eval folders yet.

Artifact by EOD: Boot log (commands that worked) + meeting contact map (names → ownership).

If install eats the 2h, stop at green-ish CI and continue Day 1.

Day 1 — Listen for the loop

Time budget: ~2h | *Mode:* meetings-heavy — listen, map, don't build

Goal: Verbal model of $x \rightarrow \text{score/decide} \rightarrow \text{log} \rightarrow \text{outcome} \rightarrow \text{learn}$.

Do today (product + agent OS merged):

- In meetings, capture: decision unit, what is logged, per-company config, what blocks rollout.
- Ask the 10 onboarding questions when natural (see §10) — especially “how do we know model M for company C is healthy?”
- Draw an 8–12 box loop on paper/whiteboard photo; don't polish.

Agent OS touch (keep tiny):

- Optional 15-min agent: /map-repo style prompt — “list train/serve/config paths” — paste into system-card; you verify tomorrow.
- Rule file optional: “read system-card before coding” — skip if no time.

Artifact by EOD: Loop sketch + raw meeting notes tagged #inventory #pain #ownership.

Day 2 — Inventory from humans, not code archaeology

Time budget: ~2h | *Mode:* meetings-heavy — listen, map, don't build

Goal: Rough model_family × company × status table.

Do today (product + agent OS merged):

- From meetings/docs: list model families (lead/churn/NBA/...) and which companies are live / pilot / blocked-on-data.
- Record folklore: last incident, who noticed, what would we miss.
- Hypothesize a pilot: 1 family + 3–5 companies (include one painful data company) — don't finalize alone.

Agent OS touch (keep tiny):

- Fill docs/onboarding/inventory.md (table ok in markdown).
- One agent pass only if you have leftover minutes: crawl configs for model-type enums — mark UNKNOWN where unsure.

Artifact by EOD: inventory.md draft v0 + pilot hypothesis (1 paragraph).

Day 3 — One thin trace (or meeting-assisted walkthrough)

Time budget: ~2h | *Mode:* meetings-heavy — listen, map, don't build

Goal: Follow one decision or request for one company as far as you can.

Do today (product + agent OS merged):

- Prefer pairing/meeting: “show me one live or staging call for company C.”
- If solo: pick one fixture/log; note features in → action/score out → what was logged.
- Write 3–5 claimed invariants (propensity? schema? allow-list? or “unknown”).

Agent OS touch (keep tiny):

- Command-shaped note for later: /trace-decision output format (you fill manually this week).
- Append trace to system-card; still no eval runner.

Artifact by EOD: One trace note + invariants list (even if half are “unknown”).

If no env access yet, a whiteboard walkthrough with an owner counts.

Day 4 — Pain list + metric hypotheses (thinking only)

Time budget: ~2h | *Mode:* meetings-heavy — listen, map, don't build

Goal: Write what “healthy” would mean before any dashboard exists.

Do today (product + agent OS merged):

- Pain list: cannot tell if... / wouldn't notice if... / data issue that looks like model issue...
- Draft metric hypotheses for the pilot family (primary + 2 slices) — label them HYPOTHESIS, not law.
- Socialize pilot hypothesis in a meeting; ask who would hate it.

Agent OS touch (keep tiny):

- Create knowledge/evals/ notes stub: “candidate metrics — not approved.”
- Create knowledge/data-readiness/ stub from meeting blockers.
- Still no skills beyond what's already stubbed.

Artifact by EOD: pain-list.md + candidate metric spec (1 page, draft).

Day 5 — Week-1 closeout / handoff to deep-work mode

Time budget: ~2h | *Mode:* meetings-heavy — listen, map, don't build

Goal: Clean package so week 2 can build without rediscovering week 1.

Do today (product + agent OS merged):

- Merge notes into: system-card.md, inventory.md, pain-list.md, pilot proposal.
- List open questions for week 2 (access gaps, missing logs, who approves metrics).
- Confirm with a stakeholder: pilot family + companies are OK to start an eval wedge.

Agent OS touch (keep tiny):

- Finalize AGENTS.md stub (hard constraints: pilot scope, don't invent KPIs, escalate promotion).
- knowledge/INDEX.md routes to stubs you actually filled.
- Explicitly schedule week-2 deep-work blocks on calendar (protect them).

Artifact by EOD: Week-1 pack: system card + inventory + pain list + pilot proposal + AGENTS.md stub + open questions.

Success for week 1 = clarity, not code. If you shipped an eval runner in week 1 you overbuilt.

Week 1 Agent OS — what you intentionally skip

- No evals/run.py, no overnight loops, no hooks, no multi-agent ablation.
- No Grafana / BI; no promotion checklist automation.
- Max tooling: AGENTS.md + a few markdown vault files + maybe one map prompt.

5. Days 6–12 — Eval harness is the product

Deep-work mode. Merged focus: ship a thin offline proof path for the pilot while adding only the commands/skills that make that path repeatable.

Day 6 — Metric spec freeze (human) + harness scaffold (agent)

Time budget: full block

Goal: Approved-enough v0 metric spec; empty runner that writes report.json shape.

Do today (product + agent OS merged):

- Turn week-1 hypotheses into a 1-page spec: primary metric, slices, fail-if hypotheses, non-goals.
- Get a quick +1 from someone who owns the pilot family (Slack is fine).
- Collect first 10 golden cases from known failures / edge companies (paths or payloads).

Agent OS touch (keep tiny):

- Agent implements evals/ scaffold to match spec (datasets/, baselines/, run → report.json).
- Add commands: /eval-spec, /eval-run (stubs OK if they call the CLI).
- Validate loop: fresh-context review of scaffold PR + attach sample report.

Artifact by EOD: evals/spec.md + runnable skeleton producing report.json + 10 case stubs.

Day 7 — Golden set growth + baseline pin

Time budget: full block

Goal: 20–30 real cases; first baseline artifact.

Do today (product + agent OS merged):

- Add cases with expected behavior + why; prefer real incidents over synthetic beauty.
- Run eval on current pilot model(s); pin baseline scores.
- Always record worst-slice, not mean only.

Agent OS touch (keep tiny):

- Skill: add-golden-case (short) — schema + escalate if label disputed.
- Command: /eval-add-case.
- Agent may propose cases from logs; you approve before merge to golden set.

Artifact by EOD: Golden set ≥20 + baselines/pilot.json + interpreted report notes.

Day 8 — Fault injection + health wedge start (parallel track begins)

Time budget: full block (split ~70% eval / 30% health)

Goal: Prove the gate can go red; start 5-minute health path.

Do today (product + agent OS merged):

- Deliberately break a model/config in a branch; show report fails; revert.
- Confirm what serving already logs (company_id, model_type, latency, status, scores).
- List missing fields that block a sanity check.

Agent OS touch (keep tiny):

- Command: /fault-inject (branch → red report → cleanup).
- Start scripts/health_check stub + /health-check command.
- Light hook optional: if api/serving touched, remind health evidence.

Artifact by EOD: Fault-injection writeup + health stub that runs for ≥1 pilot company.

Day 9 — Eval diagnostics + more cases

Time budget: full block

Goal: Layer diagnosis muscle; golden set → ~30–40.

Do today (product + agent OS merged):

- For each red/weak slice: classify data vs model vs serve vs label.
- Add cases that would have caught a past “looked good” failure.
- Document false friends (metrics that look healthy when broken).

Agent OS touch (keep tiny):

- Skill: interpret-eval-report.
- Vault: knowledge/evals/knowledge.md + rules.md (worst-slice; pin baseline; no mean-only).
- Still no overnight ablation.

Artifact by EOD: Updated golden set + eval rules in vault + diagnosis notes.

Day 10 — Wire eval into PR/CI for pilot (thin)

Time budget: full block

Goal: Regressions are harder to ignore.

Do today (product + agent OS merged):

- CI or scheduled job: evals/run --mode ci for pilot (even if only on labeled PRs).
- PR template checkbox: eval report attached when models/features change.
- Provisional thresholds — mark as hypotheses.

Agent OS touch (keep tiny):

- Cloud/background agent: “wire CI eval for pilot only; do not change metrics.”
- Fresh review + you set thresholds.
- Implement loop OK for CI plumbing; not for metric invention.

Artifact by EOD: CI/schedule path green on baseline + red on fault-inject demo.

Day 11 — Health check usable

Time budget: full block

Goal: Answer “is company C alive and non-degenerate?” in 5 minutes.

Do today (product + agent OS merged):

- Finish health_check for all pilot companies: traffic, errors, latency, flat/null scores.
- 2–3 alerts max (or manual daily glance if alerts blocked).
- Write docs/onboarding/health-5min.md runbook.

Agent OS touch (keep tiny):

- Skill: serving-sanity.
- After any serve PR: /health-check evidence required.
- Vault: required log fields in platform-core/rules.md.

Artifact by EOD: Working /health-check + runbook + sample artifacts/health/*.json.

Day 12 — Midpoint demo

Time budget: full block

Goal: Show stakeholders: we can prove a bad model and see serving sanity for pilot.

Do today (product + agent OS merged):

- 15–20 min demo: eval pass/fail + fault inject + health glance.
- Reconfirm pilot scope; collect more golden-case ideas from the room.
- Park BI/dashboard requests: CLI + JSON first.

Agent OS touch (keep tiny):

- No new skills. Freeze eval command names.
- Capture demo script in docs for reuse.

Artifact by EOD: Demo notes + agreement to proceed to promotion checklist.

6. Days 13–18 — Health finish + promotion checklist

Merge data-readiness reality with model proof. Agent OS becomes checklist automation.

Day 13 — Data readiness checklist draft

Time budget: full block

Goal: “Data fixed” becomes checkable.

Do today (product + agent OS merged):

- List checks: schema, nulls, volume, labels, leakage/forbid fields, known bad fields per company pattern.
- Sit with data-fix owners: map each major fix → which checks unlock.
- Separate sections: DATA vs MODEL readiness.

Agent OS touch (keep tiny):

- Vault: knowledge/data-readiness/rules.md + knowledge.md.
- Start /promotion-check outputting PASS/FAIL/UNKNOWN markdown.
- AGENTS.md line: data fixed ≠ done until readiness + eval baseline exist.

Artifact by EOD: promotion-checklist.md v0 (manual-friendly).

Day 14 — Model readiness + soak definition

Time budget: full block

Goal: Model side of the checklist is concrete.

Do today (product + agent OS merged):

- Model readiness: eval vs baseline, slice floors, health soak period, owner sign-off.
- Define soak: e.g. N hours/days green health for pilot company.
- Registry sketch: which model version is live per company (YAML/table).

Agent OS touch (keep tiny):

- Skills: promotion-gate, company-onboarding (sequence only).
- Command: /data-fix-to-cases — propose golden cases when data fixes land.

Artifact by EOD: Complete checklist with both sections + registry sketch.

Day 15 — Dry-run one company through /promotion-check

Time budget: full block

Goal: Find UNKNOWN gaps; turn into tickets.

Do today (product + agent OS merged):

- Run checklist on one real company (promote or dry-run).
- Every UNKNOWN becomes a ticket with owner.
- Do not invent passing scores — leave red/unknown honest.

Agent OS touch (keep tiny):

- Agent fills UNKNOWN as ticket drafts; you prioritize.
- Validate loop on any code that auto-checks data readiness.

Artifact by EOD: Filled checklist for company C + ticket list.

Day 16 — Close critical UNKNOWN gaps

Time budget: full block

Goal: Make dry-run almost fully decided by artifacts.

Do today (product + agent OS merged):

- Implement cheapest failing checks (null rates, volume floors, eval baseline exists).
- Add 5–10 golden cases unlocked by recent data fixes.
- Re-run promotion-check.

Agent OS touch (keep tiny):

- Parallel OK: agent A adds cases; agent B implements a data check — you own thresholds.
- Fresh review both PRs.

Artifact by EOD: Fewer UNKNOWNs; updated golden set.

Day 17 — Business proxy wedge decision

Time budget: full block

Goal: Either a thin weekly proxy or an explicit NO LABELS gap.

Do today (product + agent OS merged):

- If labels/joins exist for pilot: define 1 proxy; sketch weekly job (CSV/Slack OK).
- If not: write gap doc — do not fake BI.
- Keep north-star proxies out of agent freedom (approved list only).

Agent OS touch (keep tiny):

- Command stub: /degradation-report will later include proxy or gap note.
- No dashboard project kickoff.

Artifact by EOD: Proxy spec OR “no label path” doc.

Day 18 — Promotion path socialized

Time budget: full block

Goal: Team agrees checklist is the promotion language.

Do today (product + agent OS merged):

- Walk stakeholders through dry-run results.
- Adjust checklist wording; freeze v0.
- Ensure data-fix epics reference readiness checks.

Agent OS touch (keep tiny):

- Document commands for teammates in short agent-os draft.
- Avoid new loops this day — communication day.

Artifact by EOD: Frozen PromotionChecklist v0 + stakeholder +1.

7. Days 19–30 — Degradation, ownership, freeze

Day 19 — Degradation policy v0

Time budget: full block

Goal: Written rules for red/yellow.

Do today (product + agent OS merged):

- Hard block if offline primary/worst-slice regresses beyond threshold (provisional OK).
- Warn on serving sanity drift / zero traffic.
- Document who can override a block.

Agent OS touch (keep tiny):

- /degradation-report stitches eval + health + proxy/gap.
- thresholds.yaml checked in as hypotheses.

Artifact by EOD: degradation policy doc + report command output sample.

Day 20 — First overnight loop (bounded)

Time budget: full block + overnight agent

Goal: Expand golden set without changing metrics.

Do today (product + agent OS merged):

- Morning: define incident sources / log slices for case mining.
- You approve/reject every proposed golden case before merge.
- Re-run baseline after accepted cases (expect score movement — interpret, don't panic).

Agent OS touch (keep tiny):

- Overnight cloud agent: "propose golden cases from these incidents; do not change metrics."
- No overnight serving refactors.

Artifact by EOD: Accepted cases merged + rejection log (teaches the agent OS).

Day 21 — Ablation loop (judgment stays human)

Time budget: full block

Goal: Practice agentic experiment loop on pilot.

Do today (product + agent OS merged):

- Pick one lever (feature/threshold/data fix) with a clear hypothesis.
- Agent runs evals/tabulates; you decide real lift vs leakage / slice lie.
- Record decision in vault (why you shipped or didn't).

Agent OS touch (keep tiny):

- Command/skill: /ablation + agentic-ablation skill.
- Stop condition: human decision required — never auto-merge “winner.”

Artifact by EOD: Ablation table + your decision note.

Day 22 — “How we know healthy” doc

Time budget: full block

Goal: Single runbook for model M / company C.

Do today (product + agent OS merged):

- Write the operator page: commands to run, what green looks like, who to ping.
- Include offline eval, health, checklist pointers.
- Add to onboarding pack for the next hire.

Agent OS touch (keep tiny):

- Ensure /eval-run, /health-check, /promotion-check, /degradation-report are named consistently in AGENTS.md.

Artifact by EOD: Operator runbook published internally.

Day 23–24 — Harden gates & evidence culture

Time budget: 2 days

Goal: Proof is default on model/serve PRs.

Do today (product + agent OS merged):

- Every meaningful PR uses validate loop: fresh review + evidence artifact.
- Fix flaky eval/health noise so people trust the gate.
- Optionally add hooks that remind evidence paths — only if people forget.

Agent OS touch (keep tiny):

- Parallel: Bugbot/fresh reviewer + CI babysit on your PRs.
- Trim unused command stubs.

Artifact by EOD: 2+ PRs that demonstrate the full evidence path end-to-end.

Day 25–27 — Claim ownership + 60-day plan

Time budget: 3 days

Goal: You are the pilot proof owner, not a tourist.

Do today (product + agent OS merged):

- Propose owning pilot eval + degradation; co-own checklist with data.
- 60-day choice (pick one): second model-family template OR deepen business proxies OR scale readiness checks across more companies.
- Portfolio writeup: before (couldn't tell) → after (pilot can).

Agent OS touch (keep tiny):

- Finalize docs/onboarding/agent-os.md for teammates.
- Freeze skill list — delete unused drafts.
- INDEX.md complete for pilot domains.

Artifact by EOD: Ownership proposal + 60d roadmap + agent-os.md.

Day 28–30 — Buffer / polish / teach

Time budget: 3 days

Goal: Leave the system teachable and smaller than you fear.

Do today (product + agent OS merged):

- Walk a teammate through /promotion-check and /eval-run.
- Fix docs drift; archive week-1 noise.
- Retrospective: what meetings taught you that code didn't.

Agent OS touch (keep tiny):

- No new frameworks. Optional: one more overnight case expansion.
- Confirm AGENTS.md DoD matches reality.

Artifact by EOD: Teammate can run the three core commands without you; personal retro notes.

8. Canonical workflows & Agent OS reference

Workflows

A. Feature PR

spec → agent implements → fresh review → tests + /eval-run and/or /health-check → you merge

B. Is company C OK?

/promotion-check --company C → tickets for UNKNOWN → never promote on vibes

C. Did the model get worse?

/eval-run --baseline pinned → /health-check → /degradation-report → interpret worst-slice

D. Overnight (only after day ~12)

Cloud: propose golden cases; do not change metrics

Morning: you accept/reject → optional /ablation; you decide

Target layout (create gradually)

```
AGENTS.md
docs/onboarding/{system-card,inventory,pain-list,health-5min,agent-os}.md
docs/plans/observability-and-eval-foundation.md
knowledge/{INDEX,platform-core,data-readiness,evals,model-families}/
.cursor/{rules,commands,skills}/ hooks.json (late)
evals/{datasets,baselines,reports}/ + run CLI
scripts/health_check.py
```

Command map by phase

When	Commands
Week 1	/map-repo, /run-smoke (optional); manual traces
Days 6–12	/eval-spec, /eval-add-case, /eval-run, /fault-inject, /health-check
Days 13–18	/promotion-check, /data-fix-to-cases
Days 19–30	/degradation-report, /ablation

Efficiency rules

Do	Don't
Grow vault from meetings + real failures	Write 15 skills in week 1
One harness command > pretty dashboard	Dashboard-first
You spec metrics; agents implement runners	Let agents choose success metrics
Parallelize cases & CI wiring	Parallelize strategy
Week 1: notes > code	Build eval runner during meeting week

9. Repo foundation plan (team epic paste)

Shorter paste block for Linear/docs — full rationale lives in the day plan above.

Problem

Many models per company; rollouts blocked on data quality; no shared eval harness; no easy API/business metrics; no reliable degradation signal. Scaling companies without proof increases silent failure risk.

Outcomes

- Offline eval CLI → report.json with worst-slice + baseline delta (pilot family first).
- Serving health: traffic/errors/latency/sanity per company + model type.
- Promotion checklist: data readiness ≠ model readiness; both required.
- Degradation v0: block/warn on regression; fault-injection demo.

Non-goals v0

- Full BI suite; multi-agent training mesh; all model types on day one.

Phases

- 0 Inventory & pilot → 1 Eval harness → 2 Health → 3 Checklist + degradation → 4 Proxy wedge → 5 Generalize

Immediate tickets

- Inventory model×company + failure folklore
- Pilot selection + metric spec v0
- Eval runner + golden set + baseline + fault inject
- Health check script + 5-min runbook
- Promotion checklist dry-run
- CI/schedule eval gate

10. Checklists, red flags, AGENTS.md skeleton

Week checkpoints

Week	Must produce
1 (2h/day)	System card, inventory, pain list, pilot proposal, AGENTS.md stub — clarity not code
2	Eval harness v0, golden set, baseline, fault inject, health stub
3	Promotion checklist dry-run, data↔case coupling, proxy-or-gap decision
4	Degradation path, ownership + 60d plan, teammate can run core commands

Questions for meetings (week 1)

- 1. What is the production decision unit and what gets logged?
- 2. What is forbidden in learning/serving (leakage, features, labels)?
- 3. What gate promotes a model/company today (if any)?
- 4. Last bad ship — how did we catch it (or not)?
- 5. Where do offline metrics lie?
- 6. Who owns on-call / drift / retrain / data promotion?
- 7. Feature/schema contract per company?
- 8. How do we know model M for company C is healthy right now?
- 9. What should a new person not touch for two weeks?
- 10. What proof do you want on AI-assisted PRs?

Red flags

- Building eval runner or 10 skills during meeting-heavy week 1
- Fine-tuning without golden set; mean-only metrics; dashboard with no cases
- “Data fixed” with no readiness check; onboarding more companies with no degradation signal
- Can't explain last week's merged agent diff

AGENTS.md skeleton

```
# Agent instructions
## Product
Multi-tenant models (lead/churn/NBA/...). Companies roll out as data becomes ready.
## Hard constraints
- Pilot scope only unless human expands it
- Never invent business KPIs; use approved metric spec
- Data readiness ≠ model readiness; both required to promote
- Always report worst-slice, not mean only
- Attach evidence (eval report and/or health output) for model/serve PRs
- Escalate ambiguous product/threshold/promotion decisions to the human
## Before coding
Read: docs/onboarding/system-card.md, knowledge/INDEX.md, domain rules
## Commands
Week1: /map-repo /run-smoke
Later: /eval-run /eval-add-case /fault-inject /health-check
      /promotion-check /degradation-report /ablation
## Definition of done
Tests green + evidence artifact + fresh review
+ human approval on product/threshold/promotion changes
```

Bottom line

Week 1: meetings are the work — spend ~2h/day capturing the loop, inventory, pain, and pilot into a vault agents will use later. **Weeks 2–4:** build the minimum proof system (eval + health + promotion checklist) with agent OS growing only as fast as those harnesses. Clarity first; code second; dashboards last.

— End of plan —