

SPOTIO & NBA Project Playbook

Funding context, strategic fit, and a phased rollout plan for next-best-action at a capital-efficient field-sales SaaS company.

Compiled from strategy conversation — June 2026. Methods and architecture mapped to the NBA prototype (bandits, OPE, relational dataset, eval gate).

1. SPOTIO context

SPOTIO is a mature, capital-efficient niche SaaS company: founded 2014, one funding round (~\$5M Series A, 2020), ~62–68 employees, ~\$8M revenue, ~1,500 customers, slow-and-steady growth. Product: territory management, route optimization, GPS-timestamped door-knock activity tracking, pipeline, CRM sync. Mobile-first, offline-capable. AI today = DASH (LLM co-pilot, launched April 2026). Predictive ML (lead scoring, next-best-action) remains underused — the predictive layer is greenfield.

Strategic fit

- **Data fits the thesis:** years of GPS-timestamped door knocks across many customers = real relational + spatiotemporal + tabular field-sales data.
- **Greenfield ownership:** predictive NBA is underbuilt; you define the function, infra, and eval discipline.
- **Contrarian vertical from inside the leader:** field-sales / go-to-market is unsexy and money-rich; SPOTIO is a leading player.
- **Edge/offline matches edge-ML:** reps in the field need fast, cached, on-device recommendations — not live API calls mid-knock.
- **Stability suits long compounding:** a 12-year-old, profitable-ish company supports a sane, multi-year build — not hypergrowth burnout.

2. Is one small funding round a red flag?

Why it is reassuring (for stability + long-term build)

- Survived without rescue rounds, down-rounds, or runway-death layoffs — durable business.
- Founder/investor control, no desperate burn, no pivot-or-die pressure.
- Good for a role optimizing balance, Barista-FIRE, and multi-year compounding.

Legitimate yellow flags

- Modest ambition and tight R&D; budget — ML must justify itself in revenue terms, fast.
- Company likely run for cash flow and a future modest/PE exit, not moonshots.
- Remote international contractors can be exposed in cost-cutting phases.
- No equity + modest company = no wealth lottery ticket (fine if salary + reputation are the goals).

Net: the single small round should not bother you if you are optimizing for stability, data access, and a niche research reputation — not hypergrowth equity. The real watch-item is whether leadership will

fund and stay patient with the ML bet; de-risk by proving ROI early.

3. NBA prototype → SPOTIO mapping

Prototype component	SPOTIO production analog
Event store	Decision → outcome logging (append-only)
Bandits + propensity	Exploration policy + OPE-ready logs
DR / OPE gate	Promotion gate: ship only on proven lift
LightGBM $q(x,a)$	Baseline next-best-action model
Routing / TSP-P	Feed/augment existing route optimization
Ethics allow-list	Fairness, GDPR, no protected features
Relational dataset + RDL	Ambitious phase (months 12–24)
Experiment leaderboard	Institutionalized experiment tracking

Core principle: sequence **trust** → **infra** → **ambition**. Earn credibility with simple, rigorously evaluated models before relational GNNs or tabular foundation models.

4. Phased rollout at SPOTIO

Phase 0 — Land, scope, de-risk (weeks 0–4)

- Define the decision: which lead/door to target next, what action (knock / call / follow-up / skip), when — to maximize conversion or revenue per rep-hour.
- Pick one primary business metric (e.g. conversions or revenue per visit) and guardrails (adoption, fairness, no degradation).
- Audit production data: GPS knocks, dispositions, outcomes, territories, CRM data — volume, cleanliness, labels, leakage, relational structure per tenant.
- Align stakeholders on the success metric and how ROI will be measured before building.
- Build vs. augment: NBA feeds existing route optimization (rank leads → route engine sequences them); do not replace routing.

Gate: Enough clean data and leadership appetite to proceed?

Phase 1 — Offline foundation + measurement (months 1–3)

- Stand up agent-legible repo: pipeline, types, tests, AGENTS.md, one-command dev env.
- Build eval harness keyed to the business metric + promotion gate (beat baseline + clear uncertainty bar).
- Instrument decision→outcome logging with propensities from day one — before the model is good. Start light exploration so logged data supports OPE later.
- Ship dead-simple baseline NBA (heuristic or LightGBM on tabular features); grade offline against the harness. No relational GNNs yet.

Gate: Does the baseline show offline lift potential over status quo, through the gate?

Phase 2 — First online win (months 3–6)

- Randomized online experiment: holdout reps/territories get model-ranked lead lists; control keeps status quo. Measure causal lift on the primary metric.
- A clean A/B on ranked lead lists is enough first — full contextual bandit can wait.
- Treat adoption as first-class: measure adherence; work with product so recommendations land in the rep workflow (DASH / route UI). Report lift and adoption together.

Gate: Statistically credible revenue/conversion lift on the holdout — the “my model made \$X” moment.

Phase 3 — Productionize & flywheel (months 6–12)

- Harden serving for the field: precomputed/cached recommendations synced to device; monitoring, drift detection, fallback to heuristic.
- Turn logging into a data flywheel; institutionalize experiment leaderboard.
- Roll out carefully: opt-in pilot tenants, exploration caps, ethics rails.

Gate: Stable, monitored lift across multiple tenants; NBA is a relied-upon feature.

Phase 4 — Ambitious models = research bet (months 12–24)

- Relational deep learning where it beats baseline through the same gate (rep ↔ lead ↔ household ↔ territory ↔ time graph).
- Tabular foundation models for cold-start: new customer/territory/rep gets good NBA on day one via pooled transfer — key multi-tenant differentiator and publishable angle.
- Per-tenant vs. pooled modeling via relational/tabular-FM methods without data leakage.
- Publish methods (no customer data): relational NBA, cold-start tabular FM, OPE-gated deployment in production.

Gate: RDL/tabular-FM models record real lift; cold-start tenants measurably benefit.

Phase 5 — Platformize & lead (months 18–30)

- NBA becomes platform capability; experiment/leaderboard compounds improvements.
- Make the team case off proven ROI — leading buys back time (direct, don’t implement).
- Continue ~quarterly public artifacts → recognized authority at RDL × field-sales × decision-quality intersection.

Gate: Team in place; reputation and product both compounding.

5. What decides whether this works

Measure causally from day one

The most common way NBA projects die: inability to prove they caused more sales. Logging propensities + a randomized holdout from the start makes the ROI claim unarguable — and in a capital-light company, an unarguable number buys patience, budget, and eventually a team.

Adoption, not just accuracy

Build with product so reps actually follow recommendations. Measure adherence as a first-class metric. A model nobody uses produces zero lift regardless of offline metrics.

6. Week-one repo setup (summary)

- **Day 1–2:** AGENTS.md, .cursor/rules, one-command dev env (uv + Makefile), strict lint/types.
- **Day 2–3:** Profile production data; MCP DB connector; reproducible pipeline; feature allow-list.
- **Day 3–5:** Offline eval harness + promotion gate; decision/outcome logging; regression test pinning baseline.
- **End of week 1:** Gradient-boosted baseline through eval harness with one honest number — not RDL/tabular FMs yet.

Delegate to agents: pipeline scaffolding, schema exploration, feature boilerplate, tests/docs, migrations/backfills.

Keep for yourself: what the model optimizes, eval design, ship/no-ship decisions, money/privacy/security correctness, modeling approach.

Bottom line: SPOTIO gives rare raw material and ownership to execute a contrarian relational/tabular NBA bet with stability — but it will not hand you ambition or upside. Confirm data richness and publishing freedom in week one; prove causal ROI by month 6; then earn the right to the ambitious models.